

Appliances_Energy Data Set

The data set is at 10 min for about 4.5 months. The house temperature and humidity conditions were monitored with a ZigBee wireless sensor network. Each wireless node transmitted the temperature and humidity conditions around 3.3 min. Then, the wireless data was averaged for 10 minutes periods. The energy data was logged every 10 minutes with m-bus energy meters. Weather from the nearest airport weather station (Chievres Airport, Belgium) was downloaded from a public data set from Reliable Prognosis (rp5.ru), and merged together with the experimental data sets using the date and time column. Two random variables have been included in the data set for testing the regression models and to filter out non predictive attributes (parameters).

Lets get started!

Loading the data

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: %matplotlib inline
```


```
In [3]: df = pd.read_csv('energydata_complete.csv')
```

In [4]: `df.head()`

Out[4]:

	date	Appliances	lights	T1	RH_1	T2	RH_2	T3	RH_3	
0	2016-01-11 17:00:00	60	30	19.89	47.596667	19.2	44.790000	19.79	44.730000	19.00
1	2016-01-11 17:10:00	60	30	19.89	46.693333	19.2	44.722500	19.79	44.790000	19.00
2	2016-01-11 17:20:00	50	30	19.89	46.300000	19.2	44.626667	19.79	44.933333	18.92
3	2016-01-11 17:30:00	50	40	19.89	46.066667	19.2	44.590000	19.79	45.000000	18.89
4	2016-01-11 17:40:00	60	40	19.89	46.333333	19.2	44.530000	19.79	45.000000	18.89

5 rows × 29 columns



```
In [5]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19735 entries, 0 to 19734
Data columns (total 29 columns):
date                19735 non-null object
Appliances          19735 non-null int64
lights              19735 non-null int64
T1                  19735 non-null float64
RH_1                19735 non-null float64
T2                  19735 non-null float64
RH_2                19735 non-null float64
T3                  19735 non-null float64
RH_3                19735 non-null float64
T4                  19735 non-null float64
RH_4                19735 non-null float64
T5                  19735 non-null float64
RH_5                19735 non-null float64
T6                  19735 non-null float64
RH_6                19735 non-null float64
T7                  19735 non-null float64
RH_7                19735 non-null float64
T8                  19735 non-null float64
RH_8                19735 non-null float64
T9                  19735 non-null float64
RH_9                19735 non-null float64
T_out               19735 non-null float64
Press_mm_hg         19735 non-null float64
RH_out              19735 non-null float64
Windspeed           19735 non-null float64
Visibility           19735 non-null float64
Tdewpoint           19735 non-null float64
rv1                  19735 non-null float64
rv2                  19735 non-null float64
dtypes: float64(26), int64(2), object(1)
memory usage: 4.4+ MB
```

In [6]: `df.head()`

Out[6]:

	date	Appliances	lights	T1	RH_1	T2	RH_2	T3	RH_3	
0	2016-01-11 17:00:00	60	30	19.89	47.596667	19.2	44.790000	19.79	44.730000	19.00
1	2016-01-11 17:10:00	60	30	19.89	46.693333	19.2	44.722500	19.79	44.790000	19.00
2	2016-01-11 17:20:00	50	30	19.89	46.300000	19.2	44.626667	19.79	44.933333	18.92
3	2016-01-11 17:30:00	50	40	19.89	46.066667	19.2	44.590000	19.79	45.000000	18.89
4	2016-01-11 17:40:00	60	40	19.89	46.333333	19.2	44.530000	19.79	45.000000	18.89

5 rows × 29 columns



Data Cleaning as required

In [9]: `df.drop('date',axis=1, inplace=True)`

In [10]: `df.head()`

Out[10]:

	Appliances	lights	T1	RH_1	T2	RH_2	T3	RH_3	T4	
0	60	30	19.89	47.596667	19.2	44.790000	19.79	44.730000	19.000000	45.000000
1	60	30	19.89	46.693333	19.2	44.722500	19.79	44.790000	19.000000	45.000000
2	50	30	19.89	46.300000	19.2	44.626667	19.79	44.933333	18.926667	45.000000
3	50	40	19.89	46.066667	19.2	44.590000	19.79	45.000000	18.890000	45.000000
4	60	40	19.89	46.333333	19.2	44.530000	19.79	45.000000	18.890000	45.000000


5 rows × 28 columns



Exploratory Data Analysis

```
In [32]: plt.figure(figsize=(15,9))
sns.heatmap(df.isnull(), cbar=False, yticklabels=False, cmap='viridis')
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x11e6f240>
```

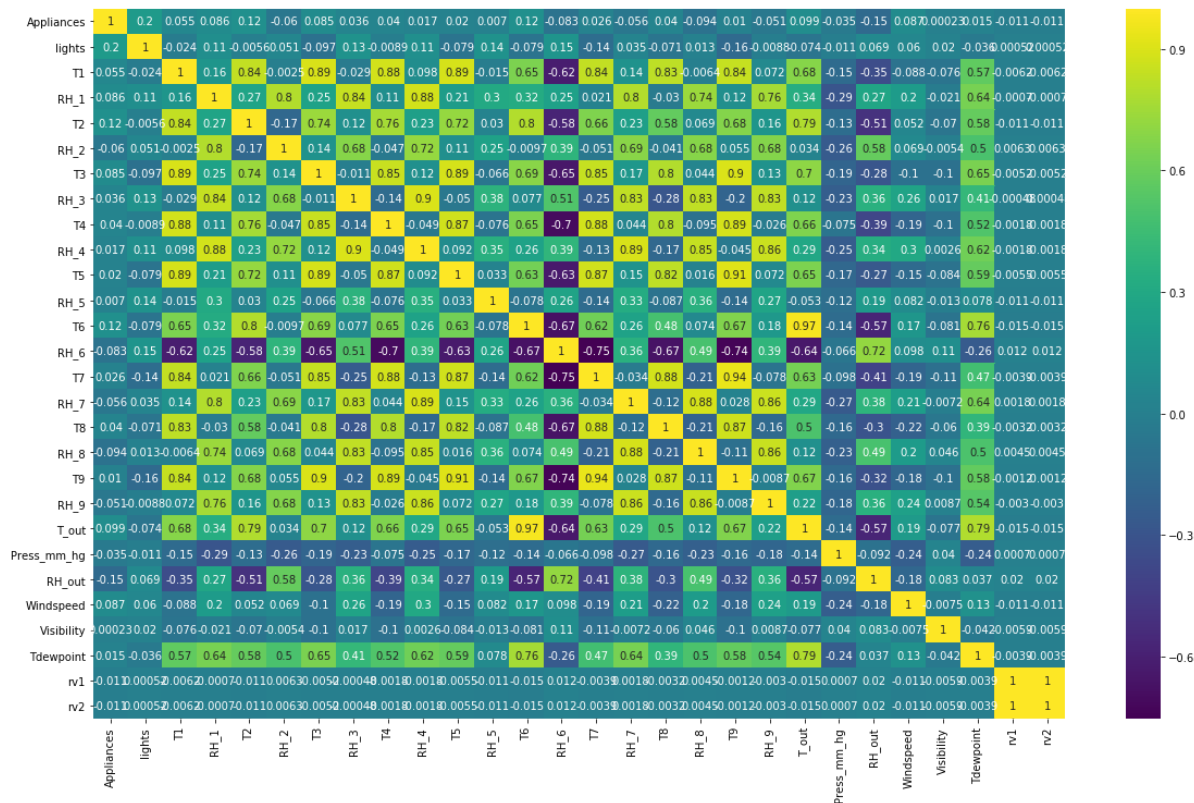


The above heatmap shows no spots, that means we do not have any null data in our dataset

Checking the correlation

```
In [30]: plt.figure(figsize=(20,12))
sns.heatmap(df.corr(), annot=True, cmap='viridis')
```

```
Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0xf90ac18>
```



Preparing Training & Test Data

```
In [11]: from sklearn.model_selection import train_test_split
```

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(df.drop('Appliances', axis
=1), df['Appliances'], test_size=0.3,
random_state=101)
```

Initializing Model & Training the same

```
In [13]: from sklearn.linear_model import LinearRegression
```

```
In [14]: lm = LinearRegression()
```

```
In [15]: lm.fit(X_train, y_train)
```

```
Out[15]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

Checking the coefficients for the trained model

```
In [16]: print(lm.coef_)
```


```
[ 2.21381857 -0.46375014 14.30189183 -16.66308548 -12.75324926
 24.29752924  4.869307    -5.11914647 -0.18411024 -1.00348894
 0.2288889   6.97473352  0.2955164   1.86226864 -1.53987426
 8.07118621 -4.57404837 -13.32788725 -0.91907821 -9.84535376
 0.18863666 -1.06362364  1.66923607  0.2074002   4.86279984
-0.03644794 -0.03644794]
```

Predicting & Representing the Values

```
In [17]: prediction = lm.predict(X_test)
```

```
In [18]: plt.scatter(y_test, prediction)
plt.xlabel('Y Values')
plt.ylabel('Predicted Values')
plt.title('Y Values Vs Predicted Values')
```

```
Out[18]: Text(0.5,1,'Y Values Vs Predicted Values')
```



Evaluating the Model

```
In [22]: from sklearn import metrics
```

```
In [23]: print("MAE: ", metrics.mean_absolute_error(y_test, prediction))
print("MSE: ", metrics.mean_squared_error(y_test, prediction))
print("RMSE: ", np.sqrt(metrics.mean_squared_error(y_test, prediction)))
```

```
MAE: 53.6273829358
MSE: 8845.48199288
RMSE: 94.0504226087
```

```
In [24]: df.columns
```

```
Out[24]: Index(['Appliances', 'lights', 'T1', 'RH_1', 'T2', 'RH_2', 'T3', 'RH_3', 'T4',
               'RH_4', 'T5', 'RH_5', 'T6', 'RH_6', 'T7', 'RH_7', 'T8', 'RH_8', 'T9',
               'RH_9', 'T_out', 'Press_mm_hg', 'RH_out', 'Windspeed', 'Visibility',
               'Tdewpoint', 'rv1', 'rv2'],
              dtype='object')
```

```
In [25]: coefficients = pd.DataFrame(lm.coef_, index=['lights', 'T1', 'RH_1', 'T2', 'RH_2', 'T3', 'RH_3', 'T4',
               'RH_4', 'T5', 'RH_5', 'T6', 'RH_6', 'T7', 'RH_7', 'T8', 'RH_8', 'T9',
               'RH_9', 'T_out', 'Press_mm_hg', 'RH_out', 'Windspeed', 'Visibility',
               'Tdewpoint', 'rv1', 'rv2'], columns=['Coefficients'])
```


In [26]: coefficients

Out[26]:

	Coefficients
lights	2.213819
T1	-0.463750
RH_1	14.301892
T2	-16.663085
RH_2	-12.753249
T3	24.297529
RH_3	4.869307
T4	-5.119146
RH_4	-0.184110
T5	-1.003489
RH_5	0.228889
T6	6.974734
RH_6	0.295516
T7	1.862269
RH_7	-1.539874
T8	8.071186
RH_8	-4.574048
T9	-13.327887
RH_9	-0.919078
T_out	-9.845354
Press_mm_hg	0.188637
RH_out	-1.063624
Windspeed	1.669236
Visibility	0.207400
Tdewpoint	4.862800
rv1	-0.036448
rv2	-0.036448


Residuals

You should have gotten a very good model with a good fit. Let's quickly explore the residuals to make sure everything was okay with our data.

Plot a histogram of the residuals and make sure it looks normally distributed. Use either seaborn distplot, or just plt.hist().

```
In [27]: sns.distplot((y_test - prediction))
```

```
Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0xc4d54a8>
```



Thus completes our project!